

The Bus Violence Dataset

The *Bus Violence* dataset is a large-scale collection of videos depicting violent and non-violent situations in public transport environments. This benchmark was gathered from multiple cameras located inside a moving bus where several people simulated violent actions, such as stealing an object from another person, fighting between passengers, etc. It contains 1,400 video clips manually annotated as having or not violent scenes, making it one of the biggest benchmarks for video violence detection in the literature.

Specifically, videos are recorded from three cameras at 25 Frames Per Second (FPS) --- two cameras located in the corners of the bus (with resolution 960x540 px) and one fisheye in the middle (1280x960 px). The clips have a minimum length of 16 frames and a maximum of 48 frames, capturing a very precise action (either violence or non-violence). The dataset is perfectly balanced, containing 700 videos of violence and 700 videos of non-violence.

The *Bus Violence* dataset is intended as a test data benchmark. However, for researchers interested in using our data also for training purposes, we provide training and test splits.

In this repository, we provide

- the 1,400 video clips, divided into two folders named Violence /NoViolence, containing clips of violent situations and non-violent situations, respectively;
- two txt files containing the names of the videos belonging to the training and test splits, respectively.